

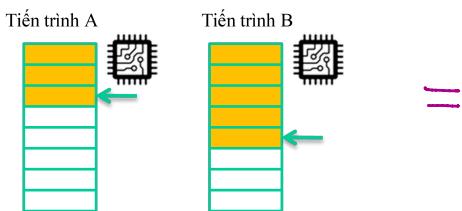
1. Processes, threads synchronization

Saturday, March 14, 2020 8:59 PM

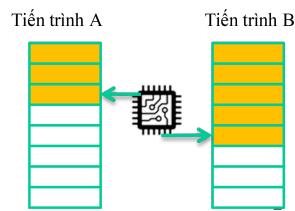
1.1 Some basic concepts

Multiple processing model

Conceptual model



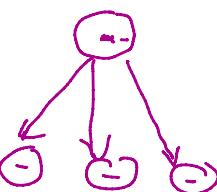
Emulation model



Parent and child processes

Unix

- A process can generate child processes
- In turn, child processes may have their own children
- All these processes can run in parallel

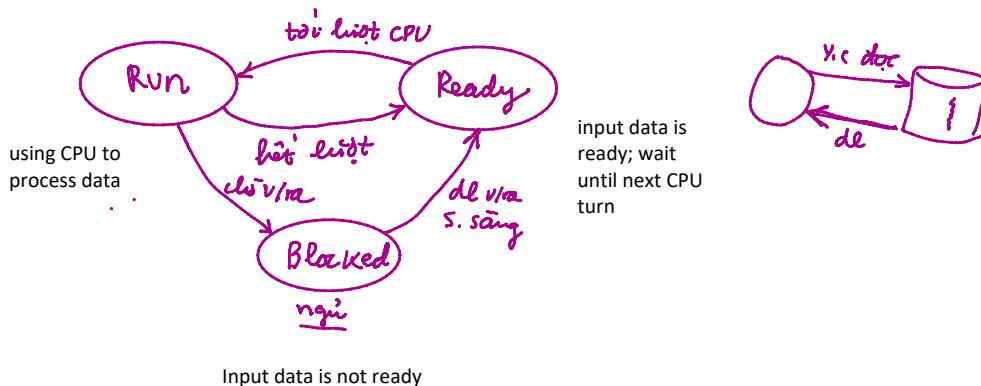


DOS

- Only one process can actually run at a time
- A process has to stall in order to load and run another process



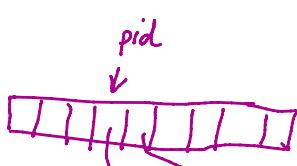
Process states



Process management

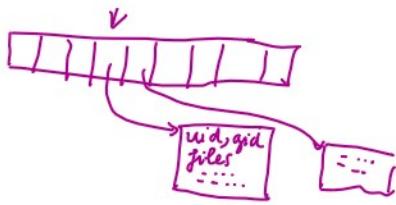
Process state table

- Process' specific data



Process state table

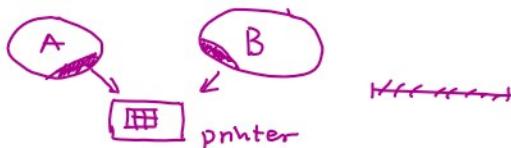
- Process' specific data
- Current memory usage
- Current files in use



Priority

- Processes with equal priority will be placed in the same queue
- Queues are managed with different priorities

Race condition



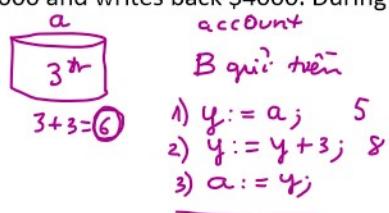
Race condition is the situation when 2 or more processes access the same exclusive shared resource concurrently and give unexpected result.

Ex 1: A and B both put data into the printer buffer, which results in data loss

Ex 2: A has \$5000 in a bank account, A withdraws \$1000 and writes back \$4000. During this process, B pays back \$3000 to A.

Tóm tắt:

- ① A₁, B₁, A₂, A₃, B₂, B₃ ②
 1) x := a; 5
 2) x := x - 2;
 ③ A₁, B₁, B₂, B₃, A₂, A₃ ④
 3) a := x;



Critical section / critical region is the part of a program if running in parallel with other codes, which access the same shared resource, may cause a race condition.

To avoid the race condition, the critical section of a program needs to be protected, not to be allowed to run in parallel with other critical codes.

1.2 Alternation method

Description

Two processes will alternately use the shared data resource in turn. After each usage, turn will be given to the other process.

Shared data

Code Int turn; // ==0 A's turn, ==1 B's turn

Process A

```
while (TRUE) {
    while (turn != 0);
    critical_section_A();
    turn = 1; // give turn to B
    non_critical_section_A();
}
```

Process B

```
while (TRUE) {
    while (turn != 1);
    critical_section_B();
    turn = 0; // give turn to A
    non_critical_section_B();}
```

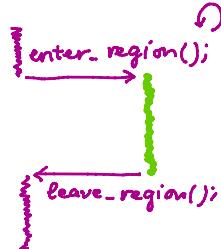
Remarks

- The number of processes is limited to 2.
- Does not separate process code from race condition handling code. Turn is hard coded into the program.
- Turn has to be alternate and has to happen even when the other process is not in need.
- Waste of CPU. Processes without resource still demand for CPU, just to check whether its turn has come.

1.3 Peterson method

Description

- Processes need to express their interest if they want to get into the critical region (to use the shared resource)
- Then, they have to call **enter_region()** to get permission to the critical region
- If passed, the process will run its critical section code
- After being done with the shared resource, the process needs to call **leave_region()** to leave the critical region and gives turn to another process



Code

Shared data

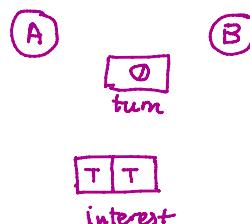
```
#define N 2; // number of processes
int turn; // ==0 A's turn, ==1 B's turn
int interest[N]; // ==T Yes, ==F No
```

Hàm enter_region() .

```
void leave_region(int pid) {
    interest[pid] = FALSE;
}
```

Hàm enter_region()

```
void enter_region(int pid) {
    int other = 1 - pid;
```



```
interest[pid] = TRUE;
turn = pid;
```

```
while (turn==pid) && (interest[other]==TRUE);
```

}

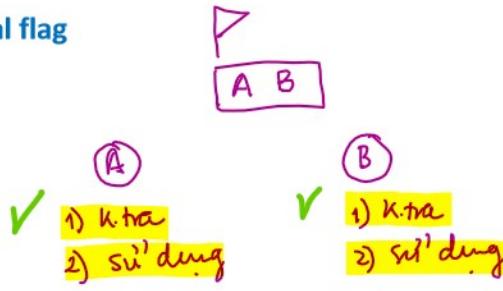
Remarks

- Race condition handling code is not hard coded into process code
For example, codes for process A and B now look like below
- process(int pid) {
 { enter_region(pid);
 critical_section(); // đoạn mã cạnh tranh
 } leave_region(pid);
 non_critical_section(); // đoạn không cạnh tranh
}
- Turn will be given only to processes that are in need of the shared resource
- The number of processes is still limited to 2
- Waste of CPU. Processes without resources still demands for CPU, just to check whether their turn have come

1.4 Test and Set Lock (TSL) using hardware

An example of using a conditional flag

```
int flag; // ==0 free, ==1 occupied  
  
L: if (flag==0) { // free  
    flag := 1; // mark as occupied  
    critical_section();  
    flag := 0; // mark as free  
}  
else goto L;
```



Test and Set Lock solution

Syntax: TSL register, flag;
equivalent to the following instructions, if executed consecutively

```
register := flag;  
flag := 1;
```

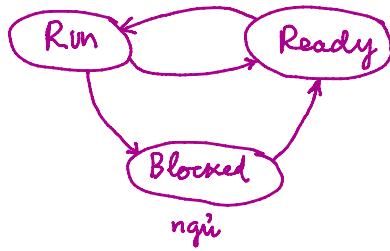
Code

```
enter_region:  
    tsl register, flag;  
    cmp register, 0; // compare register với 0  
    jnz enter_region; // Jump if Not Zero  
    ret  
  
leave_region:  
    mov flag, 0; // flag:=0 mark the resource free  
    ret
```

Remarks

- Separate race condition handling code from process code
- Turn will only be given to processes with interest of using the shared resource
- No more limitation to the number of processes
- Still, waste of CPU. Processes without resources still demand for CPU, just to check whether their turn have come

1.5 Overall remarks on the 3 above approaches

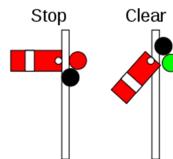


None of the above solutions turn a process into blocked state when I/O data is not available

1.4 Semaphore

Concept

- Proposed by Edsger Dijkstra in 1962
- Is a non-negative integer variable
- Its value can only be changed via 2 actions up() and down()



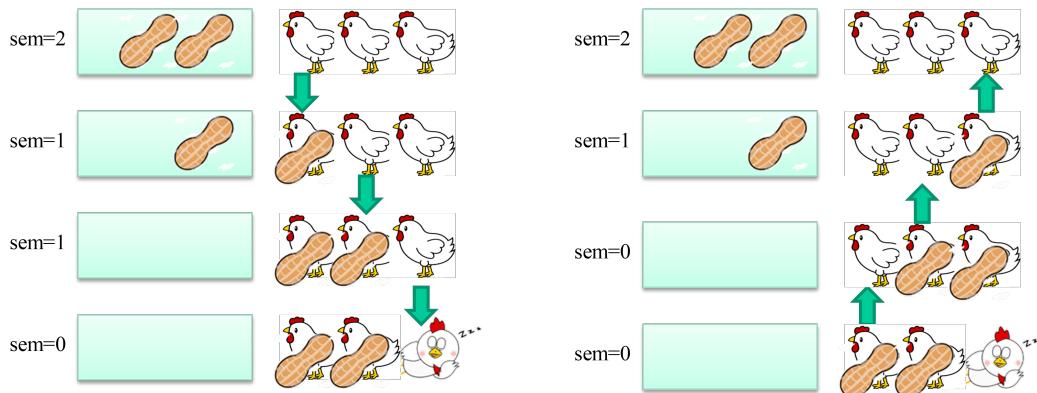
• Down(sem)

- If sem=0 the process turns into blocked state, down action is deferred
- Otherwise, sem is reduced by 1

• Up(sem)

- If sem=0, wake up a process in blocked state (sleeping on the semaphore waiting for the resource)
- Increase sem by 1

Example



1.5 Producer - Consumer problem



Người tiêu dùng Kho có n ngăn Nhà sản xuất

Consumer:

- Go to the warehouse, check for an occupied slot
- If available, get an item
- Consume the item

Producer

- Produce an item
- Go to the warehouse, check for an empty slot
- If available, put the item in

An approach

```
#define N 10 // Number of slots in the warehouse
int count; // count the number of occupied slots

void producer(void) {
    int item;
    while (TRUE) {
        produce_item(&item);
        if (count==N) sleep(); // full
        enter_item(item);
        count++;
        if (count==1) wakeup(consumer);
    }
}
```

```
void consumer(void) {
    int item;
    while (TRUE) {
        if (count==0) sleep(); // empty
        count--;
        if (count==N-1) wakeup(producer);
        consume_item(item);
    }
}
```

A problem with this approach is that race conditions may happen.

Using semaphore

Shared data

```
#define N 10 // N is the number of slots
#define int semaphore
semaphore mutex=1; // Mutual Exclusion, a binary semaphore
semaphore empty = N; // number of free slots
semaphore full; // number of occupied slots
```

Consumer code

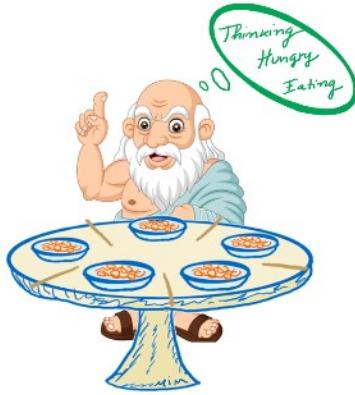
```
void consumer(void) {
    int item;
    while (TRUE) {
        down(&full); // nếu kho không có hàng thì ngủ
        down(&mutex);
        get_item(&item);
        up(&mutex);
        up(&empty); // tăng ngăn rỗng, đánh thức nsx
        consume_item(item);
    }
}
```

Producer code

```
void producer(void) {
    int item;
    while (TRUE) {
        produce_item(&item);
        down(&empty); // k tra ngăn rỗng, nếu không có thì ngủ
        down(&mutex);
        enter_item(item);
        up(&mutex);
        up(&full); // tăng ngăn có hàng, đánh thức ntd
    }
}
```

Notes: Do not perform down(&sem) within a critical section wrapped by down(&mutex) and up(&mutex) if all those semaphores relate to the same resource

1.6 Dinning philosophers problem



5 chinese philosophers sitting around a table, 5 plates of noodles in front, 5 chopsticks between the plates. A philosopher's activities:

- Thinking
- Feeling hungry, get 2 chopsticks next to him
- Eat
- Put the chopsticks down

Requirement: write code for philosophers to simulate the above activities.

An approach

```
#define N 5 // 5 philosophers
void philosopher(int i) {
    while (TRUE) {
        think();
        take_chopstick(i); T
        take_chopstick( (i+1)%N); P
        eat();
        put_chopstick(i);
        put_chopstick( (i+1)%N);
    }
}
```

Deadlock (bế tắc)

Mỗi NH_T lèn được đưa trái và đều chờ đưa phải.

take_chopsticks(i);

Using semaphore

```
// Shared data
#define N 5 // 5 philosophers
#define LEFT(i) (i-1+N)%N
#define RIGHT(i) (i+1)%N
#define THINKING 0
#define HUNGRY 1
#define EATING 2
#define typedef int semaphore

int state[N]; // state of each philosopher
```

```

semaphore s[N]; // s[i] == 1 both chopsticks are ready, ==0 not ready for philosopher i
semaphore mutex=1; // Mutual Exclusion semaphore

```

Philosopher code

```

void philosopher(int i) {
    while (TRUE) {
        think();
        take_chopsticks(i); // get 2 chopsticks or wait
        eat();
        put_chopsticks(i); // put 2 chopsticks
    }
}

```

take_chopsticks function

```

void take_chopsticks(int i) {
    down(&mutex);
    state[i] = HUNGRY;
    test(i); // check if 2 chopsticks are available. If yes, set semaphore s[i] up
    up(&mutex);
    down(&s[i]); // get the 2 chopsticks if s[i] is up, sleep otherwise
}

```

mutex : ~~ban~~

$\checkmark \text{down}(\&\text{mutex}) \equiv \text{freeze}$
 $\text{up}(\&\text{mutex}) \equiv \text{free}$

test() function

```

void test(int i) {
    if (state[i] == HUNGRY &&
        state[LEFT(i)] != EATING && // left chopstick avail.
        state[RIGHT(i)] != EATING) { // right chopstick avail.
        state[i] = EATING;
        up(&s[i]); // set s[i] up if both chopsticks are avail.
    }
}

```

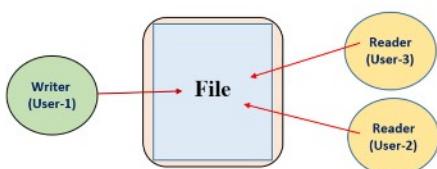
drop_chopsticks() function

```

void put_chopsticks(int i) {
    down(&mutex);
    state[i] = THINKING;
    test(LEFT(i)); // wake up the phil. on the left if sleeping
    test(RIGHT(i)); // wake up the phil on the right if sleeping
    up(&mutex);
}

```

1.7 Reader - Writer problem



	R	W
R	✓	✗
W	✗	✗

Reader:

- 1st process: get WRITE permission
- Last process: release WRITE permission
- otherwise: just read the resource

Writer:

- Get WRITE permission
- Write to the resource
- Release WRITE permission

Shared data

Tín hiệu c. trah?

- Tep : db doi ghi
- Bren rc : giam so dem t. tr. doc

```
semaphore mutex=1; // for mutual exclusion  
semaphore db = 1; // guard the write permission  
int rc=0; // reader counter
```

Reader process

```
void reader(void) {  
    while (TRUE) {  
        down(&mutex);  
        rc++; // tang so dem t. tr. doc  
        if (rc==1) down(&db); // gianh quyen ghi } c. trah  
        up(&mutex);  
        read_data(); // c. trah tep  
        down(&mutex);  
        rc--; // giam so dem t. tr. doc } c. trah  
        if (rc==0) up(&db); // tra lai quyen ghi  
        up(&mutex);  
        use_data();  
    }  
}
```

Writer process

```
void writer(void) {  
    while (TRUE) {  
        thinkup_data(); // chuan bi dl  
        down(&db); // gianh quyen ghi  
        write_data(); // c. trah tep  
        up(&db); // tra lai quyen ghi  
    }  
}
```

Question: we execute down(&db) within the scope of down(&mutex) up(&mutex). Is there a risk that all processes will be blocked waiting for each other?

Answer: no, as semaphore mutex guards for the resource rc, semaphore db guards for the file resource. These 2 semaphore are independent.

1.8 Sleeping barber problem



A barbershop with n chairs

The barber:

- Call a waiting client for the haircut, sleep if none
- If there is a sleeping client, wake him/her up
- Cut client's hair
- Back to the 1st step

Customers

- Check for any available chair, leave otherwise
- Seat and wait on a chair, call for the barber and sleep if not available
- Have the hair cut
- Go home

Using semaphore

Shared data

```
#define CHAIRS 5 // 5 chairs
```

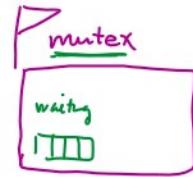
```
semaphore customers = 0; // no of customers waiting on the chairs
semaphore barbers = 0; // no of available barbers
semaphore mutex = 1; // mutual exclusion semaphore
int waiting = 0; // same value as customers
```

Barber process

```
void barber(void) {
    while (TRUE) {
        down(&customers); // sleep if no customers
        down(&mutex); // get access to 'waiting'
        waiting--;
        up(&barbers);
        up(&mutex)
        cut_hair();
    }
}
```

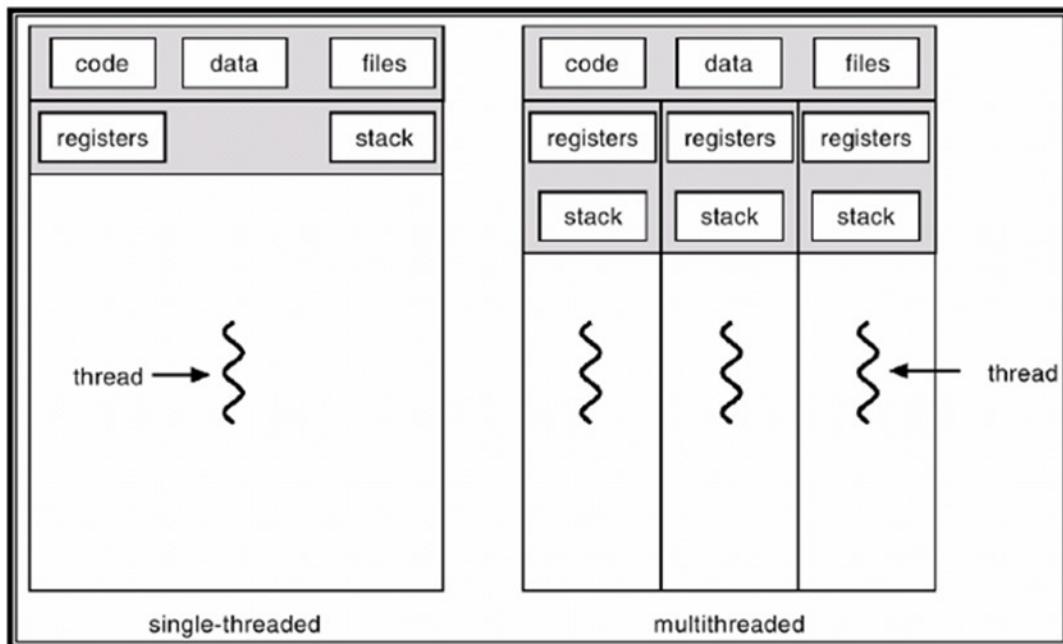
Customer process

```
void customer(void) {
    down(&mutex);
    if (waiting < CHAIRS) { // chairs available
        waiting++;
        up(&customers); // wake up a sleeping barber
        up(&mutex);
        down(&barbers); // sleep if no available barbers
        get_hair_cut();
    }
    else up(&mutex); // leave the barbershop
}
```



Threads synchronization using monitors and condition variables

Threads



A process consists of

- program code;
- process data including global variables;
- open files/resources;
- program stack containing arguments, local variables when a sub function is called;
- a set of registers that CPU is working on, including the instruction pointer that points to the next instruction.

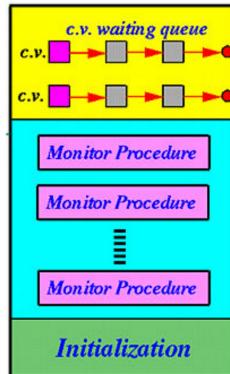
A thread is a sub environment in a process consisting of:

- Shared environment between threads: code, data segments and open resources
- Private environment:
 - o A set of register values which is specific to each thread. All threads have the same code segment but each has its own instruction pointer
 - o A stack keeping the value of sub function arguments and local variables

Monitors

A monitor is a class object with following structure. It is where concurrent programming meets object programming.

```
monitor <monitor name> {  
    shared data;  
    condition variables;  
  
    procedure init(...);  
    procedure p_1(...);  
    ...  
    procedure p_n(...);  
}
```



Monitor procedures $p_1 \dots p_n$ contain critical sections to access common resource labeled as shared data. Condition variables are used to control when to execute a critical section by using the sleep and wake up mechanism. In a monitor, only one procedure is allowed to run at a time.

Monitors are implemented at the programming language level, not at the kernel level as with semaphores. Monitors are used to synchronize threads within a process.

Example: account deposit and withdrawal (without condition variables) written in java

```
class account {  
    int balance;  
    public synchronized void deposit(int amount) {  
        int x = balance;  
        x += amount;  
        balance = x;  
    }  
    public synchronized void withdraw(int amount) {
```

```

        int y = balance;
        y -= amount;
        balance = y;
    }
};

```

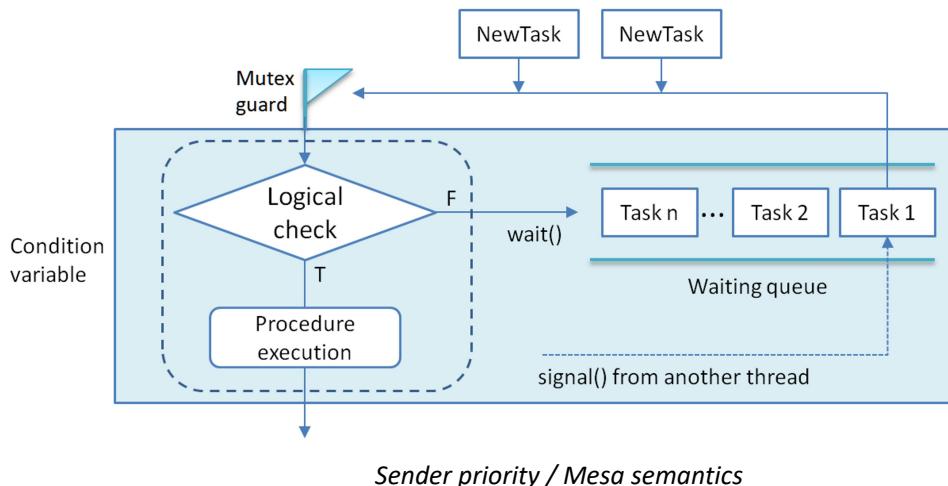
To guarantee that only one procedure is running at a moment, an automatic mutex locking/unlocking mechanism is implemented at the programming language level. The actual code should look like the one below

<pre>public synchronized void deposit(int amount) { int x; lock(this.mutex); x = balance; x += amount; balance = x; unlock(this.mutex); }</pre>	<pre>public synchronized void withdraw(int amount) { int y; lock(this.mutex); y = balance; y -= amount; balance = y; unlock(this.mutex); }</pre>
---	--

Condition variables

A condition variable is a queue structure attached to a logical condition. The queue contains threads waiting till their turn. When the logical condition is satisfied, the thread at the beginning of the queue will be woken up to run. The following functions can be used with condition variables

- `wait()` block the calling thread and add to the waiting queue
- `signal()` wake up a waiting thread (if any), otherwise the signal will be ignored



There are 2 different semantics when `signal()` is called

- Hoare semantics (receiver priority): Block the calling thread and immediately give CPU control to the woken thread. It is difficult to implement in practice.
- Mesa semantics (sender priority): Wake up a sleeping thread from the queue and move it to ready state waiting till CPU turn. The signaling thread is allowed to complete its task and exits the monitor. It is easier to implement but may cause a race condition problem as another thread may access the monitor before the woken thread. A solution is to check the logical condition again after being woken up.

Example: Producers - Consumers problem

Hoare semantics / receiver priority	Mesa semantics / sender priority
<pre>monitor ProducerConsumer { int fullSlots=0; cond empty, full; producer() { if (fullSlots==0) wait(empty); ... // fill a slot fullSlots++; signal(full); } consumer() { if (fullSlots==0) wait(full); ... // empty a slot fullSlots--; signal(empty); } }</pre>	<pre>monitor ProducerConsumer { int fullSlots=0; cond empty, full; producer() { while (fullSlots==0) wait(empty); ... // fill a slot fullSlots++; signal(full); } consumer() { while (fullSlots==0) wait(full); ... // empty a slot fullSlots--; signal(empty); } }</pre>

Implementation in C/C++

C/C++ does not provide an automatic locking/unlocking mechanism so programmers have to do it manually.

Functions for monitor programming in C/C++

```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
int pthread_mutex_init (pthread_mutex_t *mutex, const
                        pthread_mutex_attr *attr);
int pthread_mutex_lock (pthread_mutex_t *m);
int pthread_mutex_unlock (pthread_mutex_t *m);

pthread_cond_t vc = PTHREAD_COND_INITIALIZER;
int pthread_cond_init (pthread_cond_t *vc, const pthread_cond_attr *attr);
int pthread_cond_wait (pthread_cond_t *vc, pthread_mutex_t *m);
int pthread_cond_signal (pthread_cond_t *vc);
int pthread_cond_broadcast (pthread_cond_t *vc);
```

Remarks:

- `pthread_cond_wait()` blocks the calling thread and then release mutex lock so other threads can access the monitor.
- When a woken thread continues to run, mutex lock will be set back so the thread can run its code without being intervened by other threads.

Example: Producers - Consumers

void * producer (void * arg) {	void * consumer (void * arg) {
--------------------------------	--------------------------------

<pre> int item; produce_item(&item); pthread_mutex_lock(&mutex); while (count == N) pthread_cond_wait(&empty, &mutex); enter_item(item); pthread_cond_signal(&full); pthread_mutex_unlock(&mutex); } </pre>	<pre> int item; pthread_mutex_lock(&mutex); while (count == 0) pthread_cond_wait(&full, &mutex); item = get_item(); pthread_cond_signal(&empty); pthread_mutex_unlock(&mutex); consume_item(item); } </pre>
--	--

Semaphores vs monitors

Semaphores are implemented at the kernel level, can be used for synchronization of both processes or threads.	Monitors are implemented at the programming language level, can be used for threads synchronization.
Programming with semaphores can be tricky and may cause deadlocks if a process is blocked within down(&mutex) ... up(&mutex).	Programming with monitors may look simpler. Mutex locking and unlocking are automatically inserted by the compiler.
Semaphores have their own memory to keep its value. A call to sem_post() increases the semaphore value by 1 whether or not a process is sleeping on the semaphore.	A condition variables has no memory. If there is no thread sleeping on a condition variable, then a call to pthread_cond_signal() on that condition variable will simply be ignored.

Despite the differences, the 2 methods are powerfully equivalent in solving race condition problems. It is possible to implement semaphores using monitors and vice versa.

Implement semaphores using monitors

```

monitor Semaphore () {
    int count;
    condition positive;

    procedure init(int value) {
        count = value;
    }
    procedure down() {
        count--;
        if (count < 0) positive.wait();
    }
    procedure up() {
        count++;
        positive.signal();
    }
}

```

Implement monitors using semaphores

Mesa semantics (sender priority): To run a monitor procedure, a thread has to lock mutex and then unlock it when being done. Then, the woken thread can lock mutex back to continue its run. There may be a situation where another thread locks mutex successfully before the woken thread continues

to run, thus it may cause a race condition conflict. So with Mesa semantics, the thread when being woken up has to check the logical condition again, by using while() instead of if().

Shared data

```
Semaphore mutex = new Semaphore(1);
Semaphore cond = new Semaphore(0);
```

A monitor procedure

```
mutex.down();      // lock mutex
<procedure body>
mutex.up();       // unlock mutex when done
```

cond.signal in a monitor procedure

```
if (!cond.empty()) { // if there are still thread sleeping on cond
    cond.up();      // wake a thread up
}
```

cond.wait in a monitor procedure

```
mutex.up();        // before going to sleep, mutex must be realsed
cond.down();       // sleep on cond
mutex.down();     // lock mutex back, before the thread continues
```

Hoare semantics (receiver priority): A new semaphore, waitings, is introduced to control threads waiting to get access again to the monitor to complete their task. The executing thread still keeps mutex in the locked state and gives control to a waiting thread (sleeping on waitings). Mutex is only unlocked when there are no more threads sleeping on semaphore waitings. After that, other threads may lock mutex and get access to the monitor.

Shared data

```
Semaphore mutex = new Semaphore(1);
Semaphore waitings = new Semaphore(0);
Semaphore cond = new Semaphore(0);
```

A monitor procedure

```
mutex.down();
<procedure body>
if (!waitings.empty()) // if there are threads to resume its run
    waitings.up();    // wake it up, but keep mutex in locked state
else mutex.up();      // only unlock mutex when waiting threads are done
```

cond.signal in a monitor procedure

```
if (!cond.empty()) { // if there are threads sleeping on cond
    cond.up();      // wake it up
    waitings.down(); // the executing thread then sleeps, waiting for turn to get access to the
                     // monitor again
}
```

cond.wait in a monitor procedure

```
if (!waitings.empty()) // if there are threads to resume its run
    waitings.up();    // wake it up, but keep mutex in locked state
else mutex.up();      // only unlock mutex when waiting threads are done
cond.down();          // sleeping on cond
```